

# Assignment: “Calculate Primes”: (Fork/Join Practice)

Avraham Leff

November 10, 2022

## COM 2545 Assignment

### 1 Assignment-Specific Packaging

The general packaging is unchanged from the basic “Homework Requirements” (see slides from first lecture and “Homework Policies for COM 2545” on Piazza).

This assignment’s “DIR” **must be named** *PrimesFJ*.

You will implement **multiple Java classes** (see below).

Your write-up file **must be named** *PrimesFJ.pdf*.

### 2 Motivation

The assignments has multiple purposes:

- Give you practice using the JDK’s “fork/join” APIs to implement the parallel decomposition algorithm technique discussed in lecture.
- Encourage you to think about the benefits and limitations of thread-based parallelization.
- Encourage you to think more deeply about the notion of “order of growth” and problem size.

You are *encouraged* to research how fork/join works in general and the JDK implementation in particular. You are *encouraged* to use similar examples as a template and **forbidden** to view code that implements this specific problem (assuming that it exists).

You are encouraged to learn more about [prime numbers](#), but I suggest that you don't spend too much time with the topic because it won't help that much with the code ☺.

Note: the assignment is “about”

- Getting experience with “self-education”, as you research how to translate the concepts you learned in lecture into concrete API calls.
- Getting “software engineering” experience as you experiment with your implementation, and produce a useful summary of your results.

### 3 Requirements

The problem you'll be solving is: “[determine the number of primes below a given bound](#)”. The API for this problem is specified by `PrimeCalculator.nPrimesInRange`.

Your task is to solve this problem in multiple ways:

- Straightforward, “naive”, fashion: `SerialPrimes`. Do not use clever algorithmic techniques such as [Sieve of Eratosthenes](#)! We're only using this class as a baseline measurement.
- `TwoThreadPrimes`: create exactly two threads, partition the problem space between them, then use the `SerialPrimes` algorithm to solve its part of the problem. The implementation then combines the results of its two threads.
- `PrimesFJ`: use Java's *Fork/Join* framework along the lines discussed in lecture.

#### 3.1 Non-Programming

This portion of the assignment is worth 30%, and will be graded for “completeness” and “clarity”.

Your writeup must contain the following subsections, in the specified order below.

##### 3.1.1 SerialPrimes

Your writeup must address the following questions (no more than two paragraphs, ideally one).

1. If we set the `start` parameter to a constant of 2, we can characterize “problem size” in terms of the magnitude of the `end` parameter. From this perspective, what would you expect the `SerialPrimes` *order of growth* to be?
2. What is the actual order of growth?
3. Explain the observed behavior.

### 3.1.2 TwoThreadPrimes

Your writeup must address the following questions (no more than two paragraphs, ideally one).

Define algorithm performance as “*performance relative to SerialPrimes algorithm*”.

1. What is the theoretical best performance for `TwoThreadPrimes`?
2. What did you observe? Create a graph showing your results with “*increasing values of the `end` parameter (holding `start` constant at 2) on the  $x$  axis, and algorithm performance (see above) on the  $y$  axis*”
3. If there is a non-trivial discrepancy between “expected” and “actual”, provide a convincing explanation.
4. What is the observed order-of-growth?

### 3.1.3 PrimesFJ

As you know, the value that you set the *threshold* parameter is an important determinant of the performance of any *Fork/Join* algorithm. Optimize your threshold parameter to have your *Fork/Join* implementation perform as best as it can for large values of the range `2..end` where “large” means “hundreds of millions”.

Your writeup must address the following questions (no more than two paragraphs, ideally one).

Define algorithm performance as “*performance relative to SerialPrimes algorithm*”.

1. How many cores does your laptop have?

2. What is the theoretical best performance on an  $n$ -core machine?
3. What did you observe? Create a graph showing your results with “increasing values of the *end* parameter (holding *start* constant at 2) on the  $x$  axis, and algorithm performance (see above) on the  $y$  axis”
4. Pick a “very large” value of the *end* parameter (holding *start* constant at 2). Graph the performance of your *Fork/Join* algorithm on the  $y$  axis against various threshold parameters on the  $x$  axis (smaller to larger). How sensitive is your algorithm to changes in the threshold?
5. If there is a non-trivial discrepancy between “expected” and “actual”, provide a convincing explanation.
6. What is the observed order-of-growth?

Finally:

- Consider the order-of-growth of the three algorithms. Do they differ (non-trivially)? Report your results and explain them (not more than one paragraph).

### 3.2 Programming

This portion of the assignment is worth 70%, and will be graded for “correctness” and ability to meet the “Big-O” requirements at scale.

To be explicit: your algorithm implementation can use whatever set of clever techniques you choose, but **must conform to the stated design**. For example, “sequential must be sequential”, “two threads can’t be four threads”, “no caching of results”, you get the idea.

Please review the general requirements for a programming assignment! I’ve tried to reduce the chances of “mistakes” occurring through the use of a “skeleton interface”, but ultimately, **you are responsible** for ensuring that I can compile and test your code without incident.

Begin by downloading the interface and three skeleton classes from the `PrimesFJ` directory (git repository is [here](#)). Then, implement the stubbed methods.

### 3.3 SerialPrimes

This class provides a *serial (sequential)* algorithm for counting prime numbers. Your implementation should be straightforward: make sure it's correct!

### 3.4 TwoThreadPrimes

This class provides a *parallel* algorithm that uses **only two threads** to parallelize the computation.

You will be graded in terms of how closely your implementation approximates the theoretical performance improvement of two threads relative to one thread.

### 3.5 ForkJoinedPrimes

This class uses Java's *ForkJoin* framework to count the number of primes. You will be graded based on how closely your results perform relative to the theoretical maximum.